

Lab 1: Hypothesis Testing

Adam Weintraut

This is a team-based lab. Your instructor will divide you into teams of three, or possibly four students. To maximize learning, we would like all students to engage with every lab component, discussing strategy with teammates, reviewing solutions, and iterating on text and code.

The lab consists of two parts. Part one consists of foundation exercises similar to a homework. Part two is a written statistical analysis.

This lab is due before your Unit 9 live session. You will find a separate place on Gradescope to submit each part. Please submit your `.Rmd` file, as well as your `.pdf` file. This is a team submission, so one person from your team should submit the teams' work, and associate the other members of the team with the submission (this is intuitive on Gradescope.)

Contents

1 Part 1: Foundational Exercises	2
1.1 Professional Magic	2
1.2 Wrong Test, Right Data	3
1.3 Test Assumptions	4
1.3.1 World Happiness	4
1.3.2 Legislators	6
1.3.3 Wine and health	8
1.3.4 Attitudes toward the religious	10
2 Part 2: Statistical Analysis	12
2.1 Data	12
2.2 The research question	12
2.3 Guidance from political scientists	12
2.4 Report guidelines	13
General guidance	13
Introduction	13
Conceptualization and Operationalization	13
Visual Design	14
Data wrangling	14
Hypothesis testing	14
Test, results and interpretation	15

1 Part 1: Foundational Exercises

1.1 Professional Magic

Your aunt (who is a professional magician), claims to have created a pair of magical coins that share a connection to each other that makes them land in the same way. The coins are always flipped at the same time. For a given flip $i \in \{1, 2, 3, \dots\}$, let X_i be a Bernoulli random variable representing the outcome of the first coin, and let Y_i be a Bernoulli random variable representing the outcome of the second coin. You assume that each flip of the pair is independent of all other flips of the pair. For all i , you also assume that X_i and Y_i have the joint distribution given in the following table.

	$X_i = 0$	$X_i = 1$
$Y_i = 0$	$p/2$	$(1-p)/2$
$Y_i = 1$	$(1-p)/2$	$p/2$

$p \in [0, 1]$ is a parameter.

Each flip of the pair is independent of all other flips of the pair. This means that whatever happens the first time that you flip both of the coins tells you nothing about the second time that you flip both of the coins, and so on. Essentially, this is a statement about the limits of your aunt's magic.

You design the following test to evaluate your aunt's claim: You flip the coins three times, and write down that your test statistic is the sum $X_1 + Y_1 + X_2 + Y_2 + X_3 + Y_3$. That is, your test statistic is essentially the number of heads that are shown.

Your null hypothesis is that $p = 1/2$, and you plan to reject the null if your test statistic is 0 or 6.

1. What is the type 1 error rate of your test? .03125
2. What is the power of your test for the alternate hypothesis that $p = 3/4$? .105

1.2 Wrong Test, Right Data

Imagine that your organization surveys a set of customers to see how much they like your regular website, and how much they like your mobile website. Suppose that both of these preference statements are measured on 5-point Likert scales.

A Likert scale is one where a person is provided ordered categories that range from lowest to highest. You can read more about them in this seminal research design text by [Fowler](#), or this brief [overview](#). If you were to run a paired t-test using this data, what consequences would the violation of the metric scale assumption have for your interpretation of the test results? What would you propose to do to remedy this problem?

Using a paired test is correct, because our null hypothesis should be that there is no difference between how much a person likes the mobile website and the regular website. Additionally, each person answers the questions, so the data is dependent by person.

However, Likert Scales have ordinal data, and the t-test assumes metric data since it seeks to find the differences between the means of the samples. This means that we would need to make the assumption that the Likert Scale's variations in answers all mean the same thing, which they don't. This can be seen easily when we examine the difference between a 2 and a 3, versus a 4 and a 5 on the Scale. Additionally, people rank their emotions differently, so we care more about what the difference between their two rankings is.

Since we don't want to measure for the mean, and we have ordinal data, we should use Wilcoxon Paired Test instead. This will measure for paired differences between the samples instead of paired means, which is a better way to compare customer's sentiment. It also can handle signed-rank data, which is important for this "difference" calculation.

1.3 Test Assumptions

For the four following questions, your task is to evaluate the assumptions for the given test. It is not enough to say that an assumption is met or not met; instead, present your evidence in the form of background knowledge, visualizations, and numerical summaries. If you produce a histogram as part of your evaluation, be sure to consider what the most appropriate bin width is. The test that we ask you to evaluate may or may not be the most appropriate test for the scenario. Because the goal of this task is to evaluate whether the data satisfies the assumptions necessary for the test to provide meaningful results, you do not need perform the test (you may perform the test, but we will not be marking for the test results).

1.3.1 World Happiness

The file `datasets/Happiness_WHR.csv` is subsetting from the World Happiness Report, a yearly publication that uses data from the Gallup World Poll surveys. The variable life ladder is a measure of happiness, described in the FAQ as follows:

This is called the Cantril ladder: it asks respondents to think of a ladder, with the best possible life for them being a 10, and the worst possible life being a 0. They are then asked to rate their own current lives on that 0 to 10 scale. The rankings are from nationally representative samples, for the years 2018-2020.

This tries to make it so that this is a metric, but it still isn't.

You would like to know whether people in countries with high GDP per capita (higher than the mean) are more happy or less happy than people in countries with low GDP (lower than the mean).

List all assumptions for a two-sample t-test. Then evaluate each assumption, presenting evidence based on your background knowledge, visualizations, and numerical summaries.

iid paired or unpaired We should use an unpaired test, we have split into "high gdp" and "low gdp", based on the mean. Since these sets are disjoint, we

metric data Although their preface aims to make it metric, I believe this data is ordinal because it is a ranking. Using a t-test is to test the mean. normal distribution $n < 30 \rightarrow$ must be normalish $n > 30 \rightarrow$ CLT

Mann-Whitney Test: compares the medians of the two samples, more applicable for ranking.

```
happy <- read.csv('./datasets/happiness_WHR.csv')
summary(happy)
```

```
## Country.name      year      Life.Ladder      Log.GDP.per.capita
## Length:239      Min.   :2019      Min.   :2.375      Min.   : 6.966
## Class :character 1st Qu.:2019      1st Qu.:4.971      1st Qu.: 8.827
## Mode  :character Median :2019      Median :5.768      Median : 9.669
##                Mean   :2019      Mean   :5.678      Mean   : 9.584
##                3rd Qu.:2020      3rd Qu.:6.428      3rd Qu.:10.527
##                Max.    :2020      Max.    :7.889      Max.    :11.648
##                NA's     :13
## Social.support    Healthy.life.expectancy.at.birth Freedom.to.make.life.choices
```

```
## Min.      :0.4200    Min.      :48.70                Min.      :0.3850
## 1st Qu.:0.7590    1st Qu.:62.00                1st Qu.:0.7360
## Median :0.8480    Median :67.20                Median :0.8220
## Mean   :0.8256    Mean   :65.84                Mean   :0.8038
## 3rd Qu.:0.9175    3rd Qu.:70.45                3rd Qu.:0.8910
## Max.    :0.9830    Max.    :77.10                Max.    :0.9700
##                                     NA's     :8                NA's     :2
## Generosity      Perceptions.of.corruption Positive.affect Negative.affect
## Min.      :-0.28900 Min.      :0.0700                Min.      :0.3220 Min.      :0.0830
## 1st Qu.: -0.11900 1st Qu.:0.6470                1st Qu.:0.6450 1st Qu.:0.2250
## Median : -0.04500 Median :0.7780                Median :0.7330 Median :0.2830
## Mean   : -0.01524 Mean   :0.7168                Mean   :0.7149 Mean   :0.2891
## 3rd Qu.: 0.07400 3rd Qu.:0.8480                3rd Qu.:0.7900 3rd Qu.:0.3440
## Max.    : 0.56100 Max.    :0.9630                Max.    :0.8910 Max.    :0.5320
## NA's     :14      NA's     :14                NA's     :2      NA's     :2
```

```
non_null_gdp <- !is.na(happy$Log.GDP.per.capita)
```

```
# find mean --
## is our mean_gdp calculated accurately? log(mean) or mean(log) ?
mean_log_gdp <- mean(happy$Log.GDP.per.capita[non_null_gdp])
print(paste0("mean_log_gdp: ", mean_log_gdp))
```

```
## [1] "mean_log_gdp: 9.58383628318584"
```

```
# create flag for high gdp
happy$high_gdp_f <- happy$Log.GDP.per.capita > mean_log_gdp

# high gdp sample size - 121
high_gdp_sample_size <- sum(!is.na(happy$Log.GDP.per.capita[happy$high_gdp_f==T]))
print(paste0("high gdp sample size: ", high_gdp_sample_size))
```

```
## [1] "high gdp sample size: 121"
```

```
# low gdp sample size - 105
low_gdp_sample_size <- sum(!is.na(happy$Log.GDP.per.capita[happy$high_gdp_f==F]))
print(paste0("low gdp sample size: ", low_gdp_sample_size))
```

```
## [1] "low gdp sample size: 105"
```

```
## is our mean_gdp calculated accurately?
# For 2 sampled t-test, we need to assume
# - iid
# - disjoint samples based on mean gdp cutoff
# - continuous
# - gdp, log, seems continuous
# - normally distributed
# - n < 30 : need normal distribution
# - n > 30 : CLT
# - parametric or non-parametric?
# - paired or unpaired?
```

```
happy$Life.Ladder
```

```
## [1] 2.375 4.995 4.745 6.086 5.488 7.234 7.195 5.173 7.098 5.114 5.821 6.772
## [13] 4.976 5.674 6.016 3.471 6.451 5.108 4.741 4.998 4.937 7.109 4.251 5.942
## [25] 5.144 6.350 4.609 5.213 6.998 5.626 6.137 7.693 6.004 5.809 4.328 6.455
## [37] 6.035 4.100 7.780 6.690 4.914 5.164 4.892 7.035 4.967 5.952 6.262 4.768
## [49] 5.930 5.659 6.000 7.533 3.249 5.347 5.006 7.255 7.332 6.445 5.392 6.309
## [61] 5.908 4.453 6.272 4.619 6.425 6.106 5.685 5.197 5.970 4.024 3.512 5.121
## [73] 5.330 6.064 7.404 4.339 3.869 5.428 4.988 6.733 4.153 6.241 6.432 5.803
## [85] 5.563 5.386 5.057 4.932 4.434 4.436 5.449 7.425 7.205 6.113 5.004 4.356
## [97] 5.467 5.015 7.442 4.443 4.483 6.086 5.653 5.999 6.268 6.242 6.095 6.130
## [109] 5.441 3.268 6.561 5.489 6.241 3.447 6.378 6.243 6.665 5.035 5.903 6.457
## [121] 4.213 4.396 7.398 7.694 6.537 5.464 3.640 6.022 4.179 4.315 4.872 5.474
## [133] 4.948 4.702 6.711 7.157 6.944 6.600 6.154 5.081 5.467 4.197 3.307 2.694
## [145] 5.365 5.901 7.137 7.213 6.173 5.280 6.839 4.408 5.559 5.516 6.110 5.598
## [157] 4.377 5.241 7.025 6.151 5.771 5.709 6.508 6.260 6.897 7.515 5.168 5.354
## [169] 4.472 5.462 6.453 4.549 7.889 6.714 5.123 7.312 5.319 5.788 5.295 6.038
## [181] 7.575 4.225 4.865 4.785 7.035 7.195 6.488 5.257 6.118 4.094 6.168 4.547
## [193] 6.294 6.250 5.284 6.229 6.391 6.157 6.015 5.964 5.812 6.011 5.722 4.803
## [205] 4.431 4.451 7.504 7.257 5.503 5.054 7.290 5.080 6.139 5.768 5.495 6.560
## [217] 6.042 6.519 6.462 4.947 5.793 6.502 7.314 7.508 6.751 5.373 3.786 5.885
## [229] 4.731 4.862 4.641 5.270 6.458 6.798 7.028 6.310 4.574 4.838 3.160
```

```
afg_ss <- sum(!is.na(happy$Life.Ladder[happy$Country.name=='Afghanistan']))
```

1.3.2 Legislators

The file [datasets/legislators-current.csv](#) is taken from the [congress-legislators project](#) on Github. You would like to test whether Democratic or Republican senators are older.

List all assumptions for a Wilcoxon rank-sum test. Then evaluate each assumption, presenting evidence based on your background knowledge, visualizations, and numerical summaries.

- get all senators and their ages (from birthday)
- split up into democrats / republicans
-
- Could test on bucketing on age groups to define this
- Wilcoxon rank-sum
 - NO REQUIREMENT FOR NORMALITY, we don't care for this one.
- i.i.d.
 - CLT - test sample size for each group > 30
 - identical - distribution is same as what you test
 - independent - one sample has nothing to do with the other. you can't change their age
- ordinal
 - why is the age ordinal? is it?
 - * argument for: age is an integer value, not continuous (unless we include months, is this natural?)
 - * argument against: age is a metric and its distance is valuable.

- paired or unpaired?
 - unpaired,

```
legs <- read.csv('./datasets/legislators-current.csv')
summary(legs)
```

```
##   last_name      first_name    middle_name      suffix
## Length:538      Length:538    Length:538      Length:538
## Class :character Class :character Class :character Class :character
## Mode  :character Mode  :character Mode  :character Mode  :character
##
##
##
##   nickname      full_name      birthday      gender
## Length:538      Length:538    Length:538      Length:538
## Class :character Class :character Class :character Class :character
## Mode  :character Mode  :character Mode  :character Mode  :character
##
##
##
##   type      state      district      senate_class
## Length:538  Length:538    Min.   : 0.000  Min.   :1.00
## Class :character Class :character 1st Qu.: 3.000  1st Qu.:1.00
## Mode  :character Mode  :character Median : 6.000  Median :2.00
##                                     Mean  : 9.984  Mean  :2.01
##                                     3rd Qu.:13.000 3rd Qu.:3.00
##                                     Max.   :53.000  Max.   :3.00
##                                     NA's   :100    NA's   :438
##   party      url      address      phone
## Length:538    Length:538    Length:538    Length:538
## Class :character Class :character Class :character Class :character
## Mode  :character Mode  :character Mode  :character Mode  :character
##
##
##
##   contact_form      rss_url      twitter      facebook
## Length:538          Length:538    Length:538    Length:538
## Class :character    Class :character Class :character Class :character
## Mode  :character    Mode  :character Mode  :character Mode  :character
##
##
##
##   youtube      youtube_id      bioguide_id      thomas_id
## Length:538      Length:538      Length:538      Min.   : 91
## Class :character Class :character Class :character 1st Qu.:1635
## Mode  :character Mode  :character Mode  :character Median :1922
##                                     Mean  :1750
##                                     3rd Qu.:2126
##                                     Max.   :2296
```

```
##                                     NA's :209
## opensecrets_id      lis_id      fec_ids      cspan_id
## Length:538          Length:538    Length:538    Min. : 260
## Class :character    Class :character Class :character 1st Qu.: 45591
## Mode :character     Mode :character Mode :character Median : 79718
##                                     Mean : 543374
##                                     3rd Qu.:1003305
##                                     Max. : 9275683
##                                     NA's :156
## govtrack_id      votesmart_id      ballotpedia_id      washington_post_id
## Min. :300018      Min. : 119      Length:538      Mode:logical
## 1st Qu.:412199    1st Qu.: 22411      Class :character NA's:538
## Median :412570    Median : 52964      Mode :character
## Mean :412042      Mean : 75411
## 3rd Qu.:412772    3rd Qu.:133024
## Max. :456862      Max. :188334
##                                     NA's :62
## icpsr_id      wikipedia_id
## Min. :14066      Length:538
## 1st Qu.:21106    Class :character
## Median :21564    Mode :character
## Mean :24264
## 3rd Qu.:21972
## Max. :94659
## NA's :77
```

- Wicoxon rank-sum ordinal var i.i.d.

1.3.3 Wine and health

The dataset `wine` can be accessed by installing the `wooldridge` package.

```
install.packages("wooldridge")
library(wooldridge)
?wine
wine
```

It contains observations of variables related to wine consumption for 21 countries. You would like to use this data to test whether countries have more deaths from heart disease or from liver disease.

List all assumptions for a signed-rank test. Then evaluate each assumption, presenting evidence based on your background knowledge, visualizations, and numerical summaries.

Wicoxon signed-rank - paired - this is true, each row has the same value

- metric var
 - continuous value, death counts.
- i.i.d.
 - each country's deaths are independent to one another, but it's not id distributed since we only have 21 countries out of all the countries in the world. If we are trying to use this data to test whether countries have more deaths from heart or liver, we don't have enough data to represent the other countries.

- difference is symmetric
 - not sure.
 - continuous symmetric distribution?
 - how do we prove this without actually doing the test???
 -

```
# install.packages("wooldridge")
library(wooldridge)

?wine
wine
```

```
##      country alcohol deaths heart liver
## 1  Australia    2.5    785   211  15.3
## 2   Austria    3.9    863   167  45.6
## 3  Belg/Lux    2.9    883   131  20.7
## 4   Canada    2.4    793   191  16.4
## 5   Denmark    2.9    971   220  23.9
## 6   Finland    0.8    970   297  19.0
## 7    France    9.1    751    71  37.9
## 8   Iceland    0.8    743   211  11.2
## 9    Ireland    0.7   1000   300   6.5
## 10   Israel    0.6    834   183  13.7
## 11    Italy    7.9    775   107  42.2
## 12    Japan    1.5    680    36  23.2
## 13 Netherlands    1.8    773   167   9.2
## 14 New Zealand    1.9    916   266   7.7
## 15    Norway    0.8    806   227  12.2
## 16    Spain    6.5    724    86  36.4
## 17    Sweden    1.6    743   207  11.2
## 18 Switzerland    5.8    693   115  20.3
## 19         UK    1.3    941   285  10.3
## 20         US    1.2    926   199  22.1
## 21 West Germany    2.7    861   172  36.7
```

```
summary(wine)
```

```
##      country      alcohol      deaths      heart
## Length:21      Min.   :0.600      Min.   : 680      Min.   : 36.0
## Class :character 1st Qu.:1.200      1st Qu.: 751      1st Qu.:131.0
## Mode  :character Median :1.900      Median : 806      Median :191.0
##              Mean  :2.838      Mean   : 830      Mean   :183.3
##              3rd Qu.:2.900      3rd Qu.: 916      3rd Qu.:220.0
##              Max.   :9.100      Max.   :1000      Max.   :300.0
##      liver
## Min.   : 6.50
## 1st Qu.:11.20
## Median :19.00
## Mean   :21.03
## 3rd Qu.:23.90
## Max.   :45.60
```

1.3.4 Attitudes toward the religious

The file `datasets/GSS_religion` is a subset of data from the 2004 General Social Survey (GSS).

The variables `prottemp` and `cathtemp` are measurements of how a respondent feels towards protestants and towards Catholics, respectively. The GSS questions are phrased as follows:

I'd like to get your feelings toward groups that are in the news these days. I will use something we call the feeling thermometer, and here is how it works:

I'll read the names of a group and I'd like you to rate that group using the feeling thermometer. Ratings between 50 degrees and 100 degrees mean that you feel favorable and warm toward the group. Ratings between 0 degrees and 50 degrees mean that you don't feel favorable toward the group and that you don't care too much for that group. If we come to a group whose name you Don't recognize, you don't need to rate that group. Just tell me and we'll move on to the next one. If you do recognize the name, but you don't feel particularly warm or cold toward the group, you would rate the group at the 50 degree mark.

How would you rate this group using the thermometer?

You would like to test whether the US population feels more positive towards Protestants or towards Catholics.

List all assumptions for a paired t-test. Then evaluate each assumption, presenting evidence based on your background knowledge, visualizations, and numerical summaries.

paired t-test? - should be paired, each person has values for each (or gives the metrics) - same subject for each scenario

metric? - argue ordinal instead. - discrete - ranking (sentiment, feeling) - similar to ladder. - how do we explain the "difference" between values? - "how much in favor" hard to quantify

i.i.d ? - conditional approval but not enough info bla bla bla - independent - we think yes? but they need to talk about how they sampled - are the people in the same household? - demographic constraints? - identically distributed - Does this represent the entire US Population? Look at the survey to see if we are accurately representing the US pop. - maybe, can't guarantee, not enough data here.

```
religion = read.csv('./datasets/GSS_religion.csv')
summary(religion)
```

```
##           X           year           id           prottemp
## Min.      : 1.0    Min.    :2004    Min.      : 4.0    Min.      : 0.00
## 1st Qu.:201.2    1st Qu.:2004    1st Qu.: 728.8    1st Qu.: 50.00
## Median :401.5    Median :2004    Median :1373.5    Median : 60.00
## Mean   :401.5    Mean   :2004    Mean   :1381.9    Mean   : 65.56
## 3rd Qu.:601.8    3rd Qu.:2004    3rd Qu.:2053.5    3rd Qu.: 85.00
## Max.    :802.0    Max.    :2004    Max.    :2808.0    Max.    :100.00
##          cathtemp
## Min.      : 0.00
## 1st Qu.: 50.00
## Median : 60.00
## Mean   : 63.16
## 3rd Qu.: 85.00
## Max.    :100.00
```

```
# possible to recognize one group?  
# did they not include them in the sample we see? - this makes sense  
# are there any rows with a null catholic and a protestant  
sum(is.na(religion$cathtemp))
```

```
## [1] 0
```

```
sum(is.na(religion$prottemp))
```

```
## [1] 0
```

2 Part 2: Statistical Analysis

The American National Election Studies (ANES) conducts surveys of voters in the United States, with a flagship survey occurring immediately before and after each presidential election. In this part, you will use the ANES data to address a question about voters in the US. Your team will conduct a statistical analysis and generate a written report in pdf format.

This is an exercise in both statistics and professional communication. It is important that your techniques are properly executed; equally important is that your writing is clear and organized, and your argument well justified.

2.1 Data

Data for the lab should be drawn from the 2020 American National Election Studies (ANES). You can access this data at <https://electionstudies.org>. This is the official site of the ANES, a project that has [been ongoing since](#) 1948, and federally funded by the National Science Foundation since 1977.

To access the data, you will need to register for an account, confirm this account, and then login. The data that you need should come from the **2020 Time Series Study**.

You will note that there are two forms of data that are available, data that is stored in a `.dta` format, and data that is stored in a `.sav` format. Both of these are proprietary data formats (`.dta` for STATA, and `.sav` for SPSS). You will need to find an appropriate library to read this data into R; we recommend that you find a package that is within the “tidyverse”.

While you’re at the ANES website, you will also want to download the codebook, because all of the variables are marked as something like, V200002 – which isn’t very descriptive without the codebook.

For a glimpse into some of the intricacies that go into the design of this study, take a look at the introduction to the codebook.

Like many modern surveys, the ANES includes survey weights, which are used to correct for situations in which members of one demographic group are more likely to respond to the survey than members of another demographic group. (The target proportions are ultimately based on US census data). The survey weights make it possible to generalize from the a population that represents people who take the survey to a population that represents the United States as a whole. These weights are beyond the scope of our class and you are not expected to utilize them. You will still be able to learn about a population model, even if applicability to the US population is limited.

2.2 The research question

Use the ANES data to address the following question:

Did Democratic voters or Republican voters experience more difficulty voting in the 2020 election?

2.3 Guidance from political scientists

Political identification in the US is a complex phenomenon that is the topic of a large academic literature. See `./background_literature/petrocik_2009.pdf` for some guidance about how stated political identity might not match with revealed political identity at the ballot box.

As practical guidance:

1. Is it reasonable to use the vote that someone cast to identify their party preference in this case? What if someone had so difficult a time voting that they did not cast a ballot?

2. Please treat individuals who “lean” in one direction or another as members of that party. This means that someone who “Leans Democratic” should be classified as Democrat; and someone who “Leans Republican” should be classified as a Republican.

2.4 Report guidelines

This section provides some guidance for you as you write your report. In [rubric.md](#) we provide you with specific statements of how we will evaluate your report.

General guidance

- You should knit an .Rmd file to create your pdf report.
- Your report should be no more than 3 pages in standard latex formatting (i.e. `output: pdf_document`)
- You should assume your reader is familiar with statistics, but has no special knowledge of the ANES survey.
- Follow the .Rmd template that we have created, using the prompts to guide you through the parts of an analysis. Make sure you fill in each prompt with all information requested.
- Your report should contain either a plot or a table that advances the argument.

Introduction

- Begin your report with an introduction to motivate the analysis.
- Introduce the topic area and explain why the research question is interesting.
- The introduction must “do work,” connecting the general topic to the specific techniques in the report.

Conceptualization and Operationalization

How do you get from a research question to data? First, ensure that the concepts in your question are clear.

- Who or what is a voter?
- Who is a “Republican” and who is a “Democrat”?
- What is difficulty voting?

Only after you have informed your reader of what these concepts are can you then describe how you are going to *measure* these concepts.

- What would be the best **possible** method of measuring this concept? Is this method possible? Why or why not?
- What is the best **available** method of measuring this concept? Why have you opted to use this measurement instead of other possibilities? Map the concept definitions that you have written down onto the variables that you are going to use. Describe, precisely, how the variables were generated, if they come from survey data, provide the text of the question that the respondent is reacting to, not the variable name.
- What, if any, changes have you made to the dataset from how it was provided? Why did you make those changes, how much data was affected, and what are the consequences for any estimates that you produce?

Visual Design

- Any plots or tables that you include must follow basic principles of visual design.
 - A plot/figure must have a title that is informative.
 - Variables must be labeled in plain language. As an example, `v20002` does not work for a label.
 - A plot should have a good ratio of information to ink / space on the page. Do not select a large or complicated plot when a simple table conveys the same information directly.
- Do not include any plot (or R output in general) that you do not discuss in your narrative.
- The code that makes your plot/figure should be included in your report `.Rmd` file, but should not be shown in your final report. To accomplish this, you can use an `echo=FALSE` argument in the code chunk that produces the plot/figure.

Data wrangling

To answer your research question, you will have to clean, tidy, and structure the data (A.K.A. wrangle).

- The code to wrangle data should be included with your deliverable somehow. If you choose to include it in your report `.Rmd` file, then it should not be shown in the PDF of your final report. To accomplish this, you can use an `echo=FALSE` argument for the code chunk that does the wrangling.
 - A better practice – not strictly necessary for this lab – would be to write a function that loads and cleans *all* of the data that is being used by your team for its reports. This way, a single function can be run (and evaluated by your reader) for all the loading, cleaning, and manipulating.
- While we do not want to prohibit you from using additional tools for data manipulation, you should be able to complete this lab with no more than the base `stats` library, plus `dplyr` and `ggplot2` for data manipulation and plotting. Other tools within the tidyverse are available to use, but don't feel like you have to search them out.
- You will learn more by writing your own function than you would searching for a package that does one thing for your report.

Hypothesis testing

To answer your research question, you will have to execute one of the statistical tests from the course.

- The code that executes your test *should* be shown in your report, because it makes very clear the specific test that you're conducting.
- You need to argue, from the statistical principles of the course, why the test you are conducting is the *most appropriate* way to answer the research question.
- Although you might not do this for a report at your organization, for this class please list every assumption from your test, and evaluate it (assess whether the assumption is a reasonable reflection of the natural process that generated the data).
- If you identify problems with some assumptions for your test, that does not mean that you should abandon the analysis or hide the problem. If these "limitations" exist, please describe them honestly, and provide your interpretation of the consequences for your test.
- While you can choose to display the results of your test in the report, you also *certainly* need to write about these results. This should be accomplished using [inline code chunks](#), rather than by hard-coding / hard-writing output into your written report. An example of this is included in `lab_1_example_solution.Rmd`.

Test, results and interpretation

Please discuss whether any statistically significant results that you find are of *practical significance*. There are many ways to do this, but the best will provide your reader enough context to understand any measured differences in a scale appropriate to your variables. Explain the main takeaway of your analysis and how it relates to the broader context you identified in the introduction.